

# Using 'kubectl' command to interact with cluster components:

Below are some frequently used options while working with **kubectl**.

```
$ kubectl apply / create ... to create a resource
$ kubectl run ... to start a POD from an image
$ kubectl explain ... documentation of Kubernetes cluster resources
$ kubectl delete ... delete resource
$ kubectl get ... list resources
$ kubectl describe ... detailed resource information, valuable while troubleshooting
$ kubectl exec ... like docker exec command, to execute a command inside a container
$ kubectl logs ... view logs on a container
```

Along with these options we can also use some additional switches to get the output formatted.

**wide** – output additional information

**YAML** – YAML formatted API object

**JSON** – JSON formatted API object

**dry-run=<option>** – Option must be "none", "server", or "client". If client strategy, only print the object that would be sent, without sending it. If server strategy, submit server-side request without persisting the resource.

Example:

```
$ kubectl apply -f deployment.yaml --dry-run=server
$ kubectl apply -f deployment.yaml --dry-run=client ...validating
syntax in deployment.yaml file.
$ kubectl create deployment nginx --image=nginx --dry-run=client -o
yaml > deployment.yaml .... Create a manifest from an imperative kubectl
command
```

kubectl	[command]	[type]	[name]	[flags]
kubectl	get	pods	pod1	--output=yaml
kubectl	create	deployment	nginx	--image=nginx

In order to get the difference between a running Kubernetes object in kube cluster against the same being updated in a manifest file.

In below example, deployment object specified in deployment.yaml will be validated against the same deployment object deployed and running in Kubernetes cluster.

```
$ kubectl diff -f deployment.yaml
```

First, let's try to get information about the cluster

```
$ kubectl cluster-info
```

Kubernetes control plane is running at https://192.168.33.10:6443

CoreDNS is running at https://192.168.33.10:6443/api/v1/namespaces/kube-system/services/kube-dns:dns/proxy

```
$ kubectl get nodes
```

NAME	STATUS	ROLES	AGE	VERSION
kube-master	Ready	master	21m	v1.9.0
kubenode-1	Ready	<none>	2m	v1.9.0
kubenode-2	Ready	<none>	39s	v1.9.0

```
$ kubectl get nodes -o wide
```

NAME	STATUS	ROLES	AGE	VERSION	INTERNAL-IP	EXTERNAL-IP	OS-IMAGE	KERNEL-VERSION	CONTAINER-RUNTIME
controlplane	Ready	control-plane,master	19h	v1.21.1	192.168.33.10	<none>	CentOS Linux 8	4.18.0-305.3.1.el8.x86_64	docker://20.10.7
kubenode1	NotReady	<none>	19h	v1.21.1	192.168.33.11	<none>	CentOS Linux 8	4.18.0-305.3.1.el8.x86_64	docker://20.10.7

Here we can see the additional nodes which have hostname as 'kubenode-1' and 'kubenode-2'.

Now that we are done with our Master and Node setup let's start with getting the PODs ready,

There are system PODs deployed into the cluster that we can see by running command,

TO get information on how many clusters that we can manage using our `kubectl` utility, we can use the command,

```
$ kubectl config get-context
```

Getting information on POD objects.

```
$ kubectl get pods --all-namespaces
```

Or,

```
$ kubectl get pods --namespace kube-system
```

We can check the status of PODs by adding `--watch` switch,

```
$ kubectl get pods --all-namespaces --watch
```

To get some details about how client and API Server interacts, we can use the verbosity option as status below,

```
$ kubectl get pods --watch -v 6
```

To kill the `watch` or `proxy` process and come out, we use the key combination `fg` and `ctrl+c`.

To get a detailed output of all resources in the cluster run below command,

```
$ kubectl get all --all-namespaces
```

NAMESPACE	NAME	READY	STATUS	RESTARTS	AGE
kube-system	pod/coredns-558bd4d5db-6dglq	1/1	Running	1	20h
kube-system	pod/coredns-558bd4d5db-cwd5q	1/1	Running	1	20h
kube-system	pod/etcd-controlplane	1/1	Running	1	20h
kube-system	pod/kube-apiserver-controlplane	1/1	Running	1	20h
kube-system	pod/kube-controller-manager-controlplane	1/1	Running	1	20h
kube-system	pod/kube-proxy-k825c	1/1	Running	1	20h
kube-system	pod/kube-proxy-wmtrf	1/1	Running	0	20h
kube-system	pod/kube-scheduler-controlplane	1/1	Running	1	20h
kube-system	pod/weave-net-qchx1	2/2	Running	2	20h
kube-system	pod/weave-net-w95qs	1/2	CrashLoopBackOff	8	20h

NAMESPACE	NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
default	service/kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	20h
kube-system	service/kube-dns	ClusterIP	10.96.0.10	<none>	53/UDP, 53/TCP, 9153/TCP	20h

NAMESPACE	NAME	DESIRED	CURRENT	READY	UP-TO-DATE	AVAILABLE	NODE
SELECTOR	AGE						
kube-system	daemonset.apps/kube-proxy	2	2	1	2	1	
kubernetes.io/os=linux	20h						
kube-system	daemonset.apps/weave-net	2	2	1	2	1	<none>
20h							

NAMESPACE	NAME	READY	UP-TO-DATE	AVAILABLE	AGE
kube-system	deployment.apps/coredns	2/2	2	2	20h

NAMESPACE	NAME	DESIRED	CURRENT	READY	AGE
kube-system	replicaset.apps/coredns-558bd4d5db	2	2	2	20h

To get a detailed information on a resource, we can use the `explain` option on a certain resource as below,

```
$ kubectl explain pods | more
```

```
$ kubectl explain pod.spec | more
```

```
$ kubectl explain pod.spec.containers | more
```

We can also get list of all cluster resources by running command,

```
$ kubectl api-resources (Highlighted in RED are frequently used resources)
```

NAME	SHORTNAMES	APIVERSION	NAMESPACED	KIND
mutatingwebhookconfigurations		admissionregistration.k8s.io/v1	FALSE	MutatingWebhookConfiguration
validatingwebhookconfigurations		admissionregistration.k8s.io/v1	FALSE	ValidatingWebhookConfiguration
customresourcedefinitions	crd,crds	apiextensions.k8s.io/v1	FALSE	CustomResourceDefinition
apiservices		apiregistration.k8s.io/v1	FALSE	APIService
controllerrevisions		apps/v1	TRUE	ControllerRevision
daemonsets	ds	apps/v1	TRUE	DaemonSet
deployments	deploy	apps/v1	TRUE	Deployment
replicasets	rs	apps/v1	TRUE	ReplicaSet
statefulsets	sts	apps/v1	TRUE	StatefulSet
tokenreviews		authentication.k8s.io/v1	FALSE	TokenReview
localsubjectaccessreviews		authorization.k8s.io/v1	TRUE	LocalSubjectAccessReview
selfsubjectaccessreviews		authorization.k8s.io/v1	FALSE	SelfSubjectAccessReview
selfsubjectrulesreviews		authorization.k8s.io/v1	FALSE	SelfSubjectRulesReview
subjectaccessreviews		authorization.k8s.io/v1	FALSE	SubjectAccessReview
horizontalpodautoscalers	hpa	autoscaling/v1	TRUE	HorizontalPodAutoscaler
cronjobs	cj	batch/v1	TRUE	CronJob
jobs		batch/v1	TRUE	Job
certificatesigningrequests	csr	certificates.k8s.io/v1	FALSE	CertificateSigningRequest

leases		coordination.k8s.io/v1	TRUE	Lease
endpointslices		discovery.k8s.io/v1	TRUE	EndpointSlice
events	ev	events.k8s.io/v1	TRUE	Event
ingresses	ing	extensions/v1beta1	TRUE	Ingress
flowschemas		flowcontrol.apiserver.k8s.io/v1beta1	FALSE	FlowSchema
prioritylevelconfigurations		flowcontrol.apiserver.k8s.io/v1beta1	FALSE	PriorityLevelConfiguration
ingressclasses		networking.k8s.io/v1	FALSE	IngressClass
ingresses	ing	networking.k8s.io/v1	TRUE	Ingress
networkpolicies	netpol	networking.k8s.io/v1	TRUE	NetworkPolicy
runtimeclasses		node.k8s.io/v1	FALSE	RuntimeClass
poddisruptionbudgets	pdb	policy/v1	TRUE	PodDisruptionBudget
podsecuritypolicies	psp	policy/v1beta1	FALSE	PodSecurityPolicy
clusterrolebindings		rbac.authorization.k8s.io/v1	FALSE	ClusterRoleBinding
clusterroles		rbac.authorization.k8s.io/v1	FALSE	ClusterRole
rolebindings		rbac.authorization.k8s.io/v1	TRUE	RoleBinding
roles		rbac.authorization.k8s.io/v1	TRUE	Role
priorityclasses	pc	scheduling.k8s.io/v1	FALSE	PriorityClass
csidrivers		storage.k8s.io/v1	FALSE	CSIDriver
csinodes		storage.k8s.io/v1	FALSE	CSINode
storageclasses	sc	storage.k8s.io/v1	FALSE	StorageClass
volumeattachments		storage.k8s.io/v1	FALSE	VolumeAttachment
csistoragecapacities		storage.k8s.io/v1beta1	TRUE	CSIStorageCapacity
bindings		v1	TRUE	Binding
componentstatuses	cs	v1	FALSE	ComponentStatus
configmaps	cm	v1	TRUE	ConfigMap
endpoints	ep	v1	TRUE	Endpoints
events	ev	v1	TRUE	Event
limitranges	limits	v1	TRUE	LimitRange
namespaces	ns	v1	FALSE	Namespace
nodes	no	v1	FALSE	Node
persistentvolumeclaims	pvc	v1	TRUE	PersistentVolumeClaim
persistentvolumes	pv	v1	FALSE	PersistentVolume
Pods	po	v1	TRUE	Pod
podtemplates		v1	TRUE	PodTemplate
replicationcontrollers	rc	v1	TRUE	ReplicationController
resourcequotas	quota	v1	TRUE	ResourceQuota
secrets		v1	TRUE	Secret
serviceaccounts	sa	v1	TRUE	ServiceAccount

services	svc	v1	TRUE	Service
----------	-----	----	------	---------

TO get resources from a specific API group,

```
$ kubectl api-groups --api-group=<groupName> ... GroupName can be like
apps or batch ... as can be seen in the above table of resources.
```

To get detailed documentation about a resource we can use explain option with a resource.

```
$ kubectl explain pod | more
```

Or,

```
$ kubectl explain pod -recursive
```

Similar to explain more detailed information about an existing object / resource in cluster can be retrieved with option **describe**.

```
$ kubectl describe node <nodename>
```

There are two ways to work with deploying objects in Kubernetes cluster,

- **Imperative**
  - o We can deploy manage one resource / object at a time by defining requirements on command line, example,
  - o `$ kubectl create deployment webapp --image=httpd:latest`
- **Declarative**
  - o To manage multiple cluster resources / objects by defining multiple options as a code (DSC), we write a manifest file in YAML/JSON format.

## Using Imperative way:

**Generating a manifest file using the dry-run option.**

```
$ kubectl create deployment webapp \
--image=httpd:latest \
--dry-run=client -o yaml > deployment.yaml
```

As we run the above command for deployment object, we can use a similar one for creating service object and related manifest for the deployed deployment object.

```
$ kubectl expose deployment webapp \
--port=80 --target-port=8080 \
--dry-run=client -o yaml > service.yaml
```

With this command a YAML file (Manifest) for the load-balancer service will get created and the same can be deployed using `kubectl create -f service.yaml` command.

To deploy service imperatively (without Manifest) remove `--dry-run` option from command and run `kubectl expose` command to deploy service.

Once a deployment object is deployed imperatively, we can get a manifest created for the deployed object by using below command.

```
$ kubectl get deployment webapp -o deployment.yaml
```

Once objects are deployed in Kubernetes cluster, we can scale up or down the number of PODs by using scale option as shown below.

```
$ kubectl scale deployment webapp --replicas=15
```

## Starting with PODs practical (Declarative provisioning):

How to we define a API object and what goes into it, refer to the documentation..

<https://kubernetes.io/ocs/reference/kubernetes-api>

The POD is defined using a manifest file.

Here is sample manifest file.. pod.yml

```
apiVersion: v1
kind: Pod
metadata:
  name: static-web
  labels:
    role: myrole
    zone: prod
    version: v1
spec:
  containers:
  - name: web
    image: nginx
    ports:
    - name: web
      containerPort: 80
      protocol: TCP
```

Now the manifest is fed to the api-server using below command,

```
root@kubernetes-master:~# kubectl create -f pod.yml
pod "static-web" created
```

Check the status of POD using command,

```
root@kubernetes-master:~# kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
static-web	1/1	Running	0	14s

To check the POD status let's run the describe command.

```
root@kubernetes-master:~# kubectl describe pods
```

Name:	static-web
Namespace:	default
Node:	kubernetes-node-2/10.142.0.4
Start Time:	Tue, 02 Jan 2018 08:18:51 +0000
Labels:	role=myrole version=v1 zone=prod



```

Annotations:  <none>
Status:       Running
IP:           10.36.0.1
Containers:
  web:
    Container ID:
docker://8cabde7b6892ba3db454f0d73622d05659698f0ed484364c5d990b22e970fe73
    Image:      nginx
    Image ID:   docker-
pullable://nginx@sha256:cf8d5726fc897486a4f628d3b93483e3f391a76ea4897de050
0ef1f9abcd69a1
    Port:      80/TCP
    State:     Running
      Started: Tue, 02 Jan 2018 08:18:59 +0000
    Ready:     True
    Restart Count: 0
    Environment: <none>
    Mounts:
      /var/run/secrets/kubernetes.io/serviceaccount from default-token-
qtvzc (ro)
Conditions:
  Type           Status
  Initialized    True
  Ready          True
  PodScheduled   True
Volumes:
  default-token-qtvzc:
    Type:      Secret (a volume populated by a Secret)
    SecretName: default-token-qtvzc
    Optional:  false
QoS Class:     BestEffort
Node-Selectors: <none>
Tolerations:   node.kubernetes.io/not-ready:NoExecute for 300s
               node.kubernetes.io/unreachable:NoExecute for 300s

```

```

Events:
  Type    Reason             Age   From              Message
  ----    -
  Normal  Scheduled          1m    default-scheduler Successfully
assigned static-web to kubenode-2
  Normal  SuccessfulMountVolume 1m    kubelet, kubenode-2
MountVolume.SetUp succeeded for volume "default-token-qtvzc"
  Normal  Pulling            1m    kubelet, kubenode-2 pulling image
"nginx"
  Normal  Pulled              1m    kubelet, kubenode-2 Successfully
pulled image "nginx"
  Normal  Created             1m    kubelet, kubenode-2 Created
container
  Normal  Started             1m    kubelet, kubenode-2 Started
container

```

To delete a POD, we run the command,

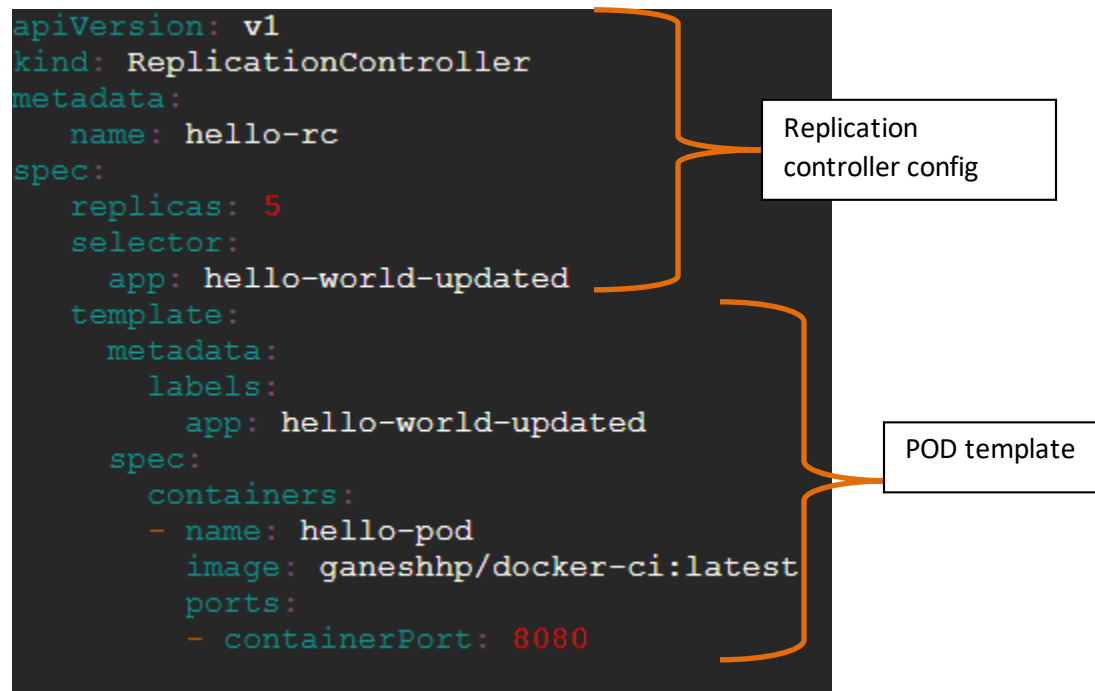
```
$ kubectl delete pods <pod-name>
```

```

root@kube-master:/home/ghpalnitkar/Kube-project# kubectl delete pods static-web
pod "static-web" deleted

```

Now let's look at **Replication Controller**, which is the right ways of applying Desired State Configuration.



So here the replication controller manifest is the actual 'Desired State Configuration' (DSC), where we state how many PODs we want to create and also supply the POD configuration that's what we have seen in earlier POD manifest.

Now try running the file one more time again using the same command,

```
$ kubectl create -f rc.yml
```

```
$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
static-web-replica-bh8rk	1/1	Running	0	23s
static-web-replica-m9ck4	1/1	Running	0	23s
static-web-replica-pkd8h	1/1	Running	0	23s
static-web-replica-qwflh	1/1	Running	0	23s
static-web-replica-t4sjs	1/1	Running	0	23s

```
$ kubectl describe pods static-web-replica-bh8rk
```

```
Name:          static-web-replica-bh8rk
Namespace:     default
Node:          kubenode-1/10.142.0.3
Start Time:    Tue, 02 Jan 2018 09:06:02 +0000
Labels:        app=static-web
Annotations:    <none>
```

```

Status:      Running
IP:          10.44.0.1
Controlled By: ReplicationController/static-web-replica
Containers:
  web:
    Container ID:
    docker://781f3009341dbc5df4beb5b54b815d39063dc2a16737cf2b7e6d71a1b06b3c27
    Image:      nginx
    Image ID:   docker-pullable://nginx@sha256:cf8d5726fc897486a4f628d3b93483e3f391a76ea4897de0500ef1f9abcd69a1
    Port:       80/TCP
    State:      Running
      Started:   Tue, 02 Jan 2018 09:06:11 +0000
    Ready:      True
    Restart Count: 0
    Environment: <none>
    Mounts:
      /var/run/secrets/kubernetes.io/serviceaccount from default-token-qtvcz (ro)
Conditions:
  Type           Status
  Initialized     True
  Ready           True
  PodScheduled    True
Volumes:
  default-token-qtvcz:
    Type:      Secret (a volume populated by a Secret)
    SecretName: default-token-qtvcz
    Optional:   false
QoS Class:     BestEffort
Node-Selectors: <none>
Tolerations:   node.kubernetes.io/not-ready:NoExecute for 300s
                node.kubernetes.io/unreachable:NoExecute for 300s
Events:
  Type     Reason             Age   From             Message
  ----     -
  Normal   Scheduled           1m    default-scheduler Successfully assigned static-web-replica-bh8rk to kubemnode-1
  Normal   SuccessfulMountVolume 1m    kubelet, kubemnode-1 MountVolume.SetUp succeeded for volume "default-token-qtvcz"
  Normal   Pulling             1m    kubelet, kubemnode-1 pulling image "nginx"
  Normal   Pulled              1m    kubelet, kubemnode-1 Successfully pulled image "nginx"
  Normal   Created             59s    kubelet, kubemnode-1 Created container
  Normal   Started             59s    kubelet, kubemnode-1 Started container

```

Now if we make any changes to the rc.yml file we can use command, to apply the changes to the cluster.

```
$ kubectl apply -f rc.yml
```

If we want to run this command with dry-run, we can use it with dry-run=server option

```
$ kubectl apply -f deployment.yaml --dry-run=server
```

Here check the age of container the earlier containers will keep on running and additional gets added.

```
root@kube-master:/home/ghpalnitkar/Kube-project# kubectl apply -f rc.yaml
Warning: kubectl apply should be used on resource created by either kubectl create --save-config or kubectl apply
replicationcontroller "static-web-replica" configured
root@kube-master:/home/ghpalnitkar/Kube-project# kubectl get pods
NAME                                READY    STATUS    RESTARTS   AGE
static-web-replica-2nxqp            1/1      Running   0           27s
static-web-replica-bh8rk            1/1      Running   0           21m
static-web-replica-jx7pk            1/1      Running   0           27s
static-web-replica-ksgsg            1/1      Running   0           27s
static-web-replica-m9ck4            1/1      Running   0           21m
static-web-replica-pkd8h            1/1      Running   0           21m
static-web-replica-qwflh            1/1      Running   0           21m
static-web-replica-t4sjs            1/1      Running   0           21m
static-web-replica-v5nzs            1/1      Running   0           27s
static-web-replica-x6dx7            1/1      Running   0           27s
root@kube-master:/home/ghpalnitkar/Kube-project#
```

Similarly if we reduce the number of container, kubernetes will make sure to destroy surplus containers and keep the one running as mentioned in the rc.yaml.

```
root@kube-master:/home/ghpalnitkar/Kube-project# kubectl apply -f rc.yaml
replicationcontroller "static-web-replica" configured
root@kube-master:/home/ghpalnitkar/Kube-project# kubectl get pods
NAME                                READY    STATUS    RESTARTS   AGE
static-web-replica-bh8rk            1/1      Running   0           25m
static-web-replica-pkd8h            1/1      Running   0           25m
static-web-replica-qwflh            1/1      Running   0           25m
static-web-replica-t4sjs            1/1      Running   0           25m
root@kube-master:/home/ghpalnitkar/Kube-project#
```

## K8S Services:

We can create a service object imperatively, by running a command as shown below.

```
$ kubectl expose rc static-web-replica --name=web-service --target-port=80 --type=NodePort
```

Here we are creating a service type object 'web-service' to and exposing the replication cluster port 80 of each pod to the service on port 80.

The type of the service that is used is NodePort service.

But running the service using command line is not always the option. We use a YAML file to define the service and tag the service to a replication controller which in turn runs PODs.

Below is such YAML file for Service.


```
apiVersion: v1
kind: Service
metadata:
  name: hello-service
  labels:
    app: hello-world-updated
spec:
  type: NodePort
  ports:
  - port: 8080
    nodePort: 30000
    protocol: TCP
  selector:
    app: hello-world-updated
```

```
$ kubectl describe service
```

```
$ kubectl describe ep <same-as-service-name>
```

## Kubernetes Deployment:

Deployments manage Replica sets and Replica sets manage PODs.



**Replication controllers** are replaced by **Replica sets** in deployment object.

```
apiVersion: extensions/v1beta1
kind: Deployment
metadata:
  name: hello-deploy
spec:
  replicas: 5
  minReadySeconds: 2
  strategy:
    type: RollingUpdate
    rollingUpdate:
      maxUnavailable: 1
      maxSurge: 1
  template:
    metadata:
      labels:
        app: hello-world-updated
    spec:
      containers:
      - name: hello-pod
        image: ganeshhp/maven-petclinic-project:latest
        ports:
        - containerPort: 8080
```

In order to allow PODs to be created on kubernetes master, use below command,

```
$ kubectl taint nodes --all node-role.kubernetes.io/master-
```

## Horizontal POD Autoscaler (HPA):

To AutoScale existing deployment object in cluster.

```
$ kubectl autoscale deployment.apps/wordpress-deploy --min=2 --max=10
```

```
$ kubectl autoscale deployment.apps/wordpress-deploy --max=10 --cpu-percent=80
```

```
$ kubectl autoscale rc.apache=rc/hello-apache --max=10 --cpu-percent=80
```

